

Class 08 Mini Project

Natalie Ogg A91030809

NOTE: I somehow messed it up a bit in the beginning so I use `diagnosis1` instead of `diagnosis` sometimes, but it's the same thing.

Today we are applying machine learning to breast cancer biopsy data from fine needle aspiration (FNA).

First I put the .csv file into the class08 file on my computer. Then I call it up and rename it:

```
wisc.df <- read.csv("WisconsinCancer.csv", row.names = 1)
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean
842302	0.11840	0.27760	0.3001	0.14710
842517	0.08474	0.07864	0.0869	0.07017
84300903	0.10960	0.15990	0.1974	0.12790
84348301	0.14250	0.28390	0.2414	0.10520
84358402	0.10030	0.13280	0.1980	0.10430
843786	0.12780	0.17000	0.1578	0.08089

	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419	0.07871	1.0950	0.9053	8.589
842517	0.1812	0.05667	0.5435	0.7339	3.398
84300903	0.2069	0.05999	0.7456	0.7869	4.585

84348301	0.2597		0.09744	0.4956	1.1560	3.445
84358402	0.1809		0.05883	0.7572	0.7813	5.438
843786	0.2087		0.07613	0.3345	0.8902	2.217
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	153.40	0.006399	0.04904	0.05373		0.01587
842517	74.08	0.005225	0.01308	0.01860		0.01340
84300903	94.03	0.006150	0.04006	0.03832		0.02058
84348301	27.23	0.009110	0.07458	0.05661		0.01867
84358402	94.44	0.011490	0.02461	0.05688		0.01885
843786	27.19	0.007510	0.03345	0.03672		0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst		
842302	0.03003		0.006193	25.38		17.33
842517	0.01389		0.003532	24.99		23.41
84300903	0.02250		0.004571	23.57		25.53
84348301	0.05963		0.009208	14.91		26.50
84358402	0.01756		0.005115	22.54		16.67
843786	0.02165		0.005082	15.47		23.75
	perimeter_worst	area_worst	smoothness_worst	compactness_worst		
842302	184.60	2019.0	0.1622			0.6656
842517	158.80	1956.0	0.1238			0.1866
84300903	152.50	1709.0	0.1444			0.4245
84348301	98.87	567.7	0.2098			0.8663
84358402	152.20	1575.0	0.1374			0.2050
843786	103.40	741.6	0.1791			0.5249
	concavity_worst	concave.points_worst	symmetry_worst			
842302	0.7119		0.2654			0.4601
842517	0.2416		0.1860			0.2750
84300903	0.4504		0.2430			0.3613
84348301	0.6869		0.2575			0.6638
84358402	0.4000		0.1625			0.2364
843786	0.5355		0.1741			0.3985
	fractal_dimension_worst					
842302		0.11890				
842517		0.08902				
84300903		0.08758				
84348301		0.17300				
84358402		0.07678				
843786		0.12440				

Now we want to omit the first column, which is the diagnosis.

```
wisc.data <- wisc.df[,-1]
head(wisc.data)
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
842302	17.99	10.38	122.80	1001.0	0.11840
842517	20.57	17.77	132.90	1326.0	0.08474
84300903	19.69	21.25	130.00	1203.0	0.10960
84348301	11.42	20.38	77.58	386.1	0.14250
84358402	20.29	14.34	135.10	1297.0	0.10030
843786	12.45	15.70	82.57	477.1	0.12780
	compactness_mean	concavity_mean	concave.points_mean	symmetry_mean	
842302	0.27760	0.3001	0.14710	0.2419	
842517	0.07864	0.0869	0.07017	0.1812	
84300903	0.15990	0.1974	0.12790	0.2069	
84348301	0.28390	0.2414	0.10520	0.2597	
84358402	0.13280	0.1980	0.10430	0.1809	
843786	0.17000	0.1578	0.08089	0.2087	
	fractal_dimension_mean	radius_se	texture_se	perimeter_se	area_se
842302	0.07871	1.0950	0.9053	8.589	153.40
842517	0.05667	0.5435	0.7339	3.398	74.08
84300903	0.05999	0.7456	0.7869	4.585	94.03
84348301	0.09744	0.4956	1.1560	3.445	27.23
84358402	0.05883	0.7572	0.7813	5.438	94.44
843786	0.07613	0.3345	0.8902	2.217	27.19
	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	0.006399	0.04904	0.05373	0.01587	
842517	0.005225	0.01308	0.01860	0.01340	
84300903	0.006150	0.04006	0.03832	0.02058	
84348301	0.009110	0.07458	0.05661	0.01867	
84358402	0.011490	0.02461	0.05688	0.01885	
843786	0.007510	0.03345	0.03672	0.01137	
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003	0.006193	25.38	17.33	
842517	0.01389	0.003532	24.99	23.41	
84300903	0.02250	0.004571	23.57	25.53	
84348301	0.05963	0.009208	14.91	26.50	
84358402	0.01756	0.005115	22.54	16.67	
843786	0.02165	0.005082	15.47	23.75	
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.60	2019.0	0.1622	0.6656	

842517	158.80	1956.0	0.1238	0.1866
84300903	152.50	1709.0	0.1444	0.4245
84348301	98.87	567.7	0.2098	0.8663
84358402	152.20	1575.0	0.1374	0.2050
843786	103.40	741.6	0.1791	0.5249
	concavity_worst	concave.points_worst	symmetry_worst	
842302	0.7119	0.2654	0.4601	
842517	0.2416	0.1860	0.2750	
84300903	0.4504	0.2430	0.3613	
84348301	0.6869	0.2575	0.6638	
84358402	0.4000	0.1625	0.2364	
843786	0.5355	0.1741	0.3985	
	fractal_dimension_worst			
842302	0.11890			
842517	0.08902			
84300903	0.08758			
84348301	0.17300			
84358402	0.07678			
843786	0.12440			

We are saving the diagnosis column for later, as a factor.

```
diagnosis1 <- as.factor(wisc.df$diagnosis)
```

Exploring the data!

Q1. Number of observations:

```
nrow(wisc.data)
```

```
[1] 569
```

Q2. How many malignant?

Use table to measure number of each character in the set:

```
table(wisc.df$diagnosis)
```

```
  B   M  
357 212
```

Other method, ask for sum where values equal M:

```
sum(wisc.df$diagnosis == "M")
```

```
[1] 212
```

Q3. How many variables/features in the data are suffixed with _mean?

grep returns the positions of matching variable names:

```
grep("_mean$", colnames(wisc.data))
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

Assign that vector to mean_vars for mean variable, then use length.

```
mean_vars <- grep("_mean$", colnames(wisc.data))  
length(mean_vars)
```

```
[1] 10
```

PCA

We need to scale our input data before PCA because the columns are measured in very different units with different means and variances. We set `scale=TRUE` argument to `prcomp()`.

`scale()` sets means to 0 and standard deviations to 1.

```
wisc.pr <- prcomp( wisc.data, scale=TRUE )  
  
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010

	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335

	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966

	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997

	PC29	PC30
Standard deviation	0.02736	0.01153
Proportion of Variance	0.00002	0.00000
Cumulative Proportion	1.00000	1.00000

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

0.4427 (from table above)

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?


3 PCs

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

7 PCs

Interpreting PCA results

```
biplot(wisc.pr)
```




Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

What stands out is everything! It is very difficult to understand because everything is dense and overlapping, with long names instead of just points.


Now we plot our PCA data and color by diagnosis:

```
plot(wisc.pr$x[,1], wisc.pr$x[,2], col=diagnosis1)
```



To add labels:

```
plot(wisc.pr$x[,1], wisc.pr$x[,2], col=diagnosis1, xlab = "PC1", ylab = "PC2")
```

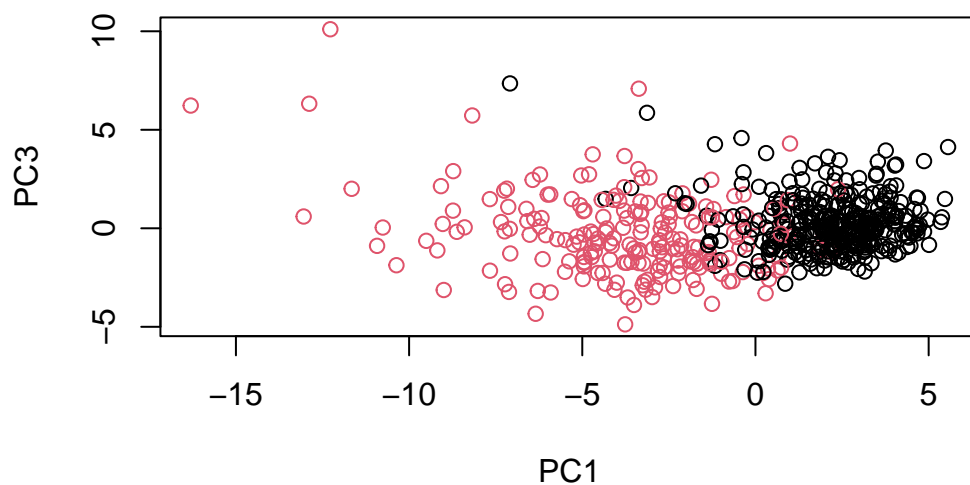



We can see that the diagnoses are starkly separated on the plot, which is notable. The idea of PCA plots here is that more similar cells will be clustered. It's a method for compressing a lot of data into something that represents the essence of the data.

You can create a point to represent a cluster of data in the PCA, for example (from class) using the original data and the PCA data for all rows to get a value for each column.

Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
plot(wisc.pr$x[, 1 ], wisc.pr$x[, 3 ], col = diagnosis1,  
     xlab = "PC1", ylab = "PC3")
```



With PC3, the points appear less clustered than in PC2, and the M and B diagnoses overlap more.

Create a data.frame for ggplot


```
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis1
```

Load the ggplot2 package

```
library(ggplot2)
```

Make a scatter plot:

```
ggplot(df) +
  aes(PC1, PC2, col=df$diagnosis) +
  geom_point()
```



Here we use SD squared to calculate the variance of each PCA component:

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```


```
[1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

Variance explained by each principal component: pve

```
pve <- pr.var / sum(pr.var)
```

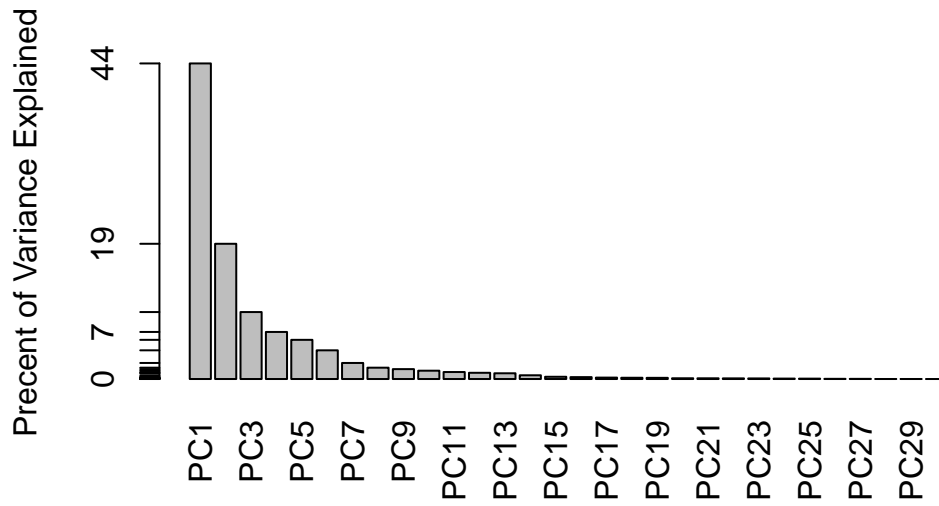
Plot variance explained for each principal component

```
plot(pve, xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



Alternative scree plot of the same data, note data driven y-axis

```
barplot(pve, ylab = "Precent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```



```
wisc.pr$rotation[,1]
```

radius_mean	texture_mean	perimeter_mean
-0.21890244	-0.10372458	-0.22753729
area_mean	smoothness_mean	compactness_mean
-0.22099499	-0.14258969	-0.23928535
concavity_mean	concave.points_mean	symmetry_mean
-0.25840048	-0.26085376	-0.13816696
fractal_dimension_mean	radius_se	texture_se
-0.06436335	-0.20597878	-0.01742803
perimeter_se	area_se	smoothness_se
-0.21132592	-0.20286964	-0.01453145
compactness_se	concavity_se	concave.points_se
-0.17039345	-0.15358979	-0.18341740
symmetry_se	fractal_dimension_se	radius_worst
-0.04249842	-0.10256832	-0.22799663
texture_worst	perimeter_worst	area_worst
-0.10446933	-0.23663968	-0.22487053
smoothness_worst	compactness_worst	concavity_worst
-0.12795256	-0.21009588	-0.22876753
concave.points_worst	symmetry_worst	fractal_dimension_worst
-0.25088597	-0.12290456	-0.13178394

Communicating PCA Results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

This is just asking for the value of `wisc.pr$rotation row concave.points_mean`

```
wisc.pr$rotation["concave.points_mean",1]
```

```
[1] -0.2608538
```

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

5. PCA 1-5 explain 80% of the variance.

```
pve
```

```
[1] 4.427203e-01 1.897118e-01 9.393163e-02 6.602135e-02 5.495768e-02  
[6] 4.024522e-02 2.250734e-02 1.588724e-02 1.389649e-02 1.168978e-02  
[11] 9.797190e-03 8.705379e-03 8.045250e-03 5.233657e-03 3.137832e-03  
[16] 2.662093e-03 1.979968e-03 1.753959e-03 1.649253e-03 1.038647e-03  
[21] 9.990965e-04 9.146468e-04 8.113613e-04 6.018336e-04 5.160424e-04  
[26] 2.725880e-04 2.300155e-04 5.297793e-05 2.496010e-05 4.434827e-06
```

```
sum(pve[1:5])
```

```
[1] 0.8473427
```

Hierarchical Clustering

Scale the `wisc.data` data using the “`scale()`” function

`data.scaled <- ____ (wisc.data)`

```
data.scaled <- scale(wisc.data)
```

Calculate the (Euclidean) distances between all pairs of observations in the new scaled dataset and assign the result to `data.dist`.

```
data.dist <- dist(data.scaled)
```


Create a hierarchical clustering model using complete linkage. Manually specify the `method` argument to `hclust()` and assign the results to `wisc.hclust`.

```
wisc.hclust <- hclust(data.dist, method = "complete")
```

Q11. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

I used `h = 20`

```
plot(wisc.hclust)
abline(h = 20, col="red", lty=2)
```



```
data.dist
hclust (*, "complete")
```

Selecting number of clusters

```
wisc.hclust.clusters <- cutree(wisc.hclust, k=4)

table(wisc.hclust.clusters, diagnosis1)
```

	diagnosis1	
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2


Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

The best match is k=4 because then the majority of the B and M diagnoses are separated into different rows (rows 1 and 3 above).

Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

```
plot(hclust(data.dist, method = "complete"))
```



Cluster Dendrogram



```
data.dist  
hclust (*, "complete")
```

```
plot(hclust(data.dist, method = "single"))
```


Cluster Dendrogram



```
data.dist  
hclust (*, "single")
```

```
plot(hclust(data.dist, method = "average"))
```


Cluster Dendrogram



data.dist
hclust (*, "average")

```
plot(hclust(data.dist, method = "ward.D2"))
```

Cluster Dendrogram




```
data.dist  
hclust (*, "ward.D2")
```

I definitely prefer the “ward.D2” clustering. It presents the cleanest groupings, and is easiest to read.

Combining methods

```
d <- dist(wisc.pr$x[,1:3])  
wisc.pr.hclust <- hclust(d, method="ward.D2")  
plot(wisc.pr.hclust)
```

Cluster Dendrogram



```
d
hclust(*, "ward.D2")
```


Generate 2 cluster groups from this hclust at the height for the number of clusters we want, here it's 2.

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
 1  2
203 366
```

Plotting with color from `grps` instead of the expert diagnosis from before, we see that we have a very similar result!

```
plot(wisc.pr$x[,1], wisc.pr$x[,2], col=grps)
```



Let's compare them:

```
table(grps, diagnosis1)
```


```

diagnosis1
grps   B   M
1    24 179
2   333  33

```

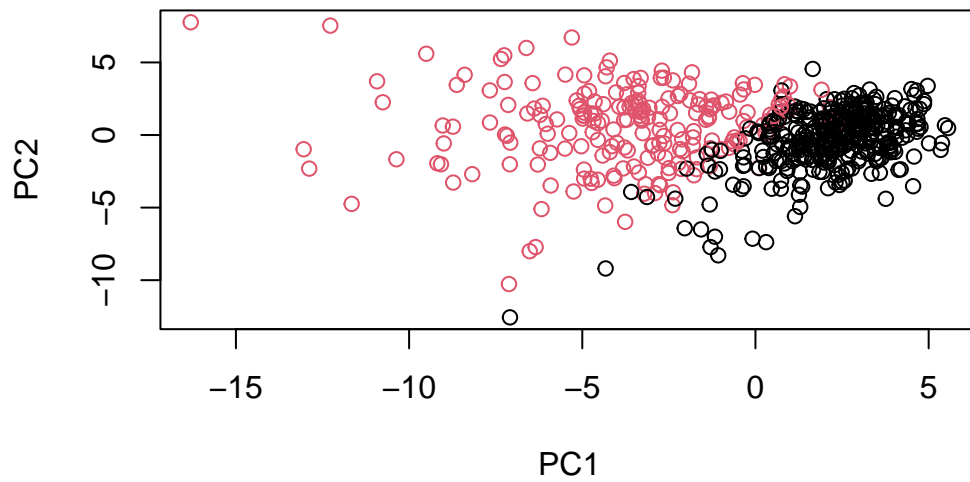
Here we color by groups:

```
plot(wisc.pr$x[,1:2], col=grps)
```



Here we color by diagnosis:

```
plot(wisc.pr$x[,1:2], col=diagnosis1)
```



Changing the colors to match: (turn grps into a factor)


```
g <- as.factor(grps)
levels(g)
```

```
[1] "1" "2"
```

```
g <- relevel(g,2)
levels(g)
```

```
[1] "2" "1"
```

```
plot(wisc.pr$x[,1:2], col=g)
```



Q15. How well does the newly created model with four clusters separate out the two diagnoses?

```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
table(wisc.pr.hclust.clusters, diagnosis1)
```


```

              diagnosis1
wisc.pr.hclust.clusters  B  M
1      24 179
2     333  33
```

THis model works fairly well because the majority of diagnoses are separated, but there are still some potential false positives and false negatives in each cluster.

```
w <- dist(wisc.pr$x[,1:7])
wisc.pr.hclust<- hclust(w, method="ward.D2")
plot(wisc.pr.hclust)
```


Cluster Dendrogram



W
hclust (*, "ward.D2")